

# Text Classification and Neural Networks

Giovanni Colavizza

Text Mining  
Amsterdam University College

March 10, 2021

# Overview

- 1 Text Classification
- 2 Logistic Regression
- 3 Evaluation
- 4 Neural Networks
- 5 Extras

## **Text Classification**

# Task definition

- We are given a **training set**  $\{X, Y\}$  of data pairs  $(x, y)$ , where  $x$  is a text document and  $y$  is the class the document belongs to.
- Each  $y \in \mathcal{Y}$ , where  $\mathcal{Y} = \{c_1, c_2, \dots, c_k\}$  are the distinct (finite and enumerable) classes we have. If  $|\mathcal{Y}| = k = 2$ , we have a binary classification task.
- Using a *learning method*, our goal is to learn a **classifier**, or a classification function  $\gamma$  that maps documents to classes:

$$\gamma : \mathcal{X} \rightarrow \mathcal{Y}$$

- The fact that we use annotated data to learn makes this a form of *supervised learning*. Note that a “document” can be anything really: words, text sequences, longer texts.

# Examples

task	$x$	$y$
language ID	text	{english, mandarin, greek, ...}
spam classification	email	{spam, not spam}
authorship attribution	text	{jk rowling, james joyce, ...}
genre classification	novel	{detective, romance, gothic, ...}
sentiment analysis	text	{postive, negative, neutral, mixed}

*Credit: David Bamman (UC Berkeley).*

# Text representation

Our text documents  $X$  can be **represented** in many ways:

- Pre-computed features (e.g., the length of the document or the average length of the words it contains).
- A selection of words (e.g., only stopwords for language detection).
- Words in isolation (so called “bag of words”, or unigram model).
- Conjunctions of words (e.g., bigrams).
- Higher-order features (e.g., PoS).
- Word embeddings.

# Logistic Regression

# Logistic regression

- Our goal is, given a document represented with a feature vector  $\mathbf{x}$  and classes  $c \in \mathcal{Y}$ , to learn a classifier discriminating the right class for  $\mathbf{x}$ :

$$\hat{p}(y = c|\mathbf{x})$$

- Let us start with a binary classifier and two classes, thus  $\mathcal{Y} = \{0, 1\}$ .
- We need to estimate  $\hat{p}(y = 1|\mathbf{x})$ , and  $\hat{p}(y = 0|\mathbf{x}) = 1 - \hat{p}(y = 1|\mathbf{x})$  will follow suit.
- Logistic regression uses two components for this: a **linear model** of the inputs and the **Sigmoid (or logistic) function**. So, it is like the perceptron but with a different classification function.



## Sigmoid (or logistic) function

- Let us consider the set of features  $x_1, x_2, \dots, x_d$  we used to represent our input document  $\mathbf{x}$ . We add  $x_0 = 1$  to model the intercept, and create a linear model with them:


$$z = \sum_{j=0}^d w_j x_j = \mathbf{w} \cdot \mathbf{x}$$

- To create a probability distribution, we pass  $z$  through the Sigmoid  $\sigma(z)$ :

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- The Sigmoid squeezes  $z$  within 0 and 1 and is always positive.

# Sigmoid (or logistic) function



**Figure 5.1** The sigmoid function  $y = \frac{1}{1+e^{-z}}$  takes a real value and maps it to the range  $[0, 1]$ . Because it is nearly linear around 0 but has a sharp slope toward the ends, it tends to squash outlier values toward 0 or 1.

*Credit: M&J, Ch. 5.*

# Logistic regression

- Applied to our binary classification task, we have that:

$$\hat{p}(y = 1|\mathbf{x}) = \sigma(z_{\mathbf{x}}) = \frac{1}{1 + e^{-\mathbf{w} \cdot \mathbf{x}}}$$
$$\hat{p}(y = 0|\mathbf{x}) = 1 - \sigma(z_{\mathbf{x}}) = \frac{e^{-\mathbf{w} \cdot \mathbf{x}}}{1 + e^{-\mathbf{w} \cdot \mathbf{x}}}$$

- Then, we just need to use a **decision boundary** to assign the class given the estimated probabilities:

$$\hat{y} = \begin{cases} 1 & \text{if } \hat{p}(y = 1|\mathbf{x}) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

- So, we have defined our data and task, and have a model.  
What do we miss?

## Logistic regression: Cross-entropy

- We need a loss function. Let us use MLE to find one.
- We have that  $p(y|\mathbf{x})$  follows a Bernoulli distribution given that we only have two discrete outcomes  $(0, 1)$ , hence:

$$p(y|\mathbf{x}) = \hat{y}^y(1 - \hat{y})^{1-y}$$

- As usual, let us move to log space and add a minus to switch to a minimization problem (note we work with a single data point  $(\mathbf{x}, y)$  for now):

$$\begin{aligned} -\log p(y|\mathbf{x}) &= -\log[\hat{y}^y(1 - \hat{y})^{1-y}] \\ &= -[y\log\hat{y} + (1 - y)\log(1 - \hat{y})] \end{aligned}$$

- Let us now plug-in the Sigmoid and call it the loss:

$$\mathcal{L}_{\mathbf{x}}(\mathbf{w}) = -[y\log\sigma(\mathbf{w}\mathbf{x}) + (1 - y)\log(1 - \sigma(\mathbf{w}\mathbf{x}))]$$

## Logistic regression: Cross-entropy

- Let us now plug-in the Sigmoid and call it the loss:

$$\mathcal{L}_x(\mathbf{w}) = -[y \log \sigma(\mathbf{w}\mathbf{x}) + (1 - y) \log(1 - \sigma(\mathbf{w}\mathbf{x}))]$$

- The loss on the whole dataset is going to be (note we are already in log space thus we can sum):

$$\mathcal{L}(\mathbf{w}) = -\frac{1}{N} \sum_{i=1}^N [y_i \log \sigma(\mathbf{w}\mathbf{x}_i) + (1 - y_i) \log(1 - \sigma(\mathbf{w}\mathbf{x}_i))]$$

- To this we can, as usual, attach regularization:

$$\mathcal{L}_{L_2}(\mathbf{w}) = \mathcal{L}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

## Logistic regression: Optimization via SGD

- The last missing bit is how to find good parameters  $\mathbf{w}$ : we can use SGD.
- It turns out that the derivative for one data point  $\mathbf{x}$  is (w.o. regularization):

$$\frac{\partial \mathcal{L}_{\mathbf{x}}(\mathbf{w})}{\partial \mathbf{w}_j} = [\sigma(\mathbf{w}\mathbf{x}) - y] \mathbf{x}_j$$

- For multiple data points, we just sum (w.o. regularization), and with this we are good to go for SGD:

$$\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}_j} = \sum_{i=1}^N [\sigma(\mathbf{w}\mathbf{x}_i) - y_i] \mathbf{x}_{ij}$$

- *Full derivation as an extra, below.*

# Notes

# Notes



## Evaluation

# Data splitting

	training	development	testing
size	80%	10%	10%
purpose	training models	model selection; hyperparameter tuning	evaluation; never look at it until the very end

*Credit: David Bamman (UC Berkeley).*

# Accuracy and baselines


- **Accuracy** is the fraction of correctly predicted data points over the total. It can be calculated on any dataset split: train, development and test. Very good starting point.
- **Baseline**: important to have one. It can be a random classifier (i.e., flip a coin for a binary classifier), or a fast and reasonable model (e.g., logistic regression with TF-IDF features).

# Precision and recall

Given a binary classifier:

- **True positive:** a data point correctly predicted to be 1.
- **True negative:** a data point correctly predicted to be 0.
- **False positive:** a data point incorrectly predicted to be 1.
- **False negative:** a data point incorrectly predicted to be 0.

# Precision and recall



How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

*Credit: Wikipedia.*

## F-measure and accuracy reloaded

- F-measure (harmonic mean of precision and recall):

$$F = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

- Accuracy:

$$A = \frac{tp + tn}{tp + tn + fp + fn}$$

# Parameters and hyperparameters

Parameters whose values are *learned*

Feature	$\beta$
the	0.01
and	0.03
bravest	1.4
love	3.1
loved	1.2
genius	0.5
<i>BIAS</i>	-0.1

Hyperparameters whose values are *chosen*


Hyperparameter	value
minimum word frequency	5
max vocab size	10000
lowercase	TRUE
regularization strength	1.0

*Credit: David Bamman (UC Berkeley).*

# Neural Networks




# A single neuron



*Credit: Andrej Karpathy via Stanford's CS231N.*

# A single neuron



*Credit: Andrej Karpathy via Stanford's CS231N.*

# Logistic regression as a neural network

Following the notation in the previous slide, we have:

- $\mathbf{x} = \langle x_0, x_1, x_2, \dots, x_d \rangle$  is our input representation.
- We aggregate the features  $\mathbf{x}$  into a linear combination using weights  $\mathbf{w}$ . We also include the bias term  $b$  into the matrix by adding an appropriate dimension fixed at 1 to  $\mathbf{x}$ , so that we can use matrix notation:



$$z = \sum_{i=0}^d w_i x_i = \mathbf{w} \cdot \mathbf{x}$$

- We pass  $z$  through an activation function, in this case the sigmoid:

$$f = \sigma(z) = \frac{1}{1 + e^{-z}}$$

- **Linear models are a single neuron.** *Question: what is the activation function for linear regression?*

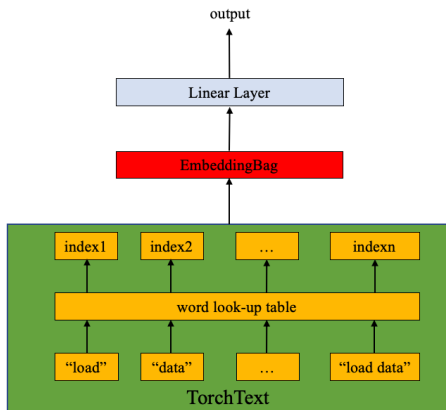
## From single layer to multi-layer



*Credit: Andrej Karpathy via Stanford's CS231N.*

## Using embeddings as features


Neural networks are **modular**: we can piece them together into advanced architectures. For example, we can use embeddings to represent our input (either training them or using pre-trained ones). *More on this in the lab.*



*Credit: TorchText.*

# Why do we need non-linearities?

Multiple layers and **non-linear functions** (such as the sigmoid) allow us to fit complex decision boundaries.



*Credit: Andrej Karpathy via Stanford's CS231N.*

# How do we train neural networks?

- Key idea: use a smart way to apply SGD, called **backpropagation**.
- Backpropagation combines using the chain rule to calculate local derivatives (called gradients) with the re-use of pre-computed operations to speed the computation up.
- *More on this in the external materials for the course.*

# Neural networks practicalities

Training neural networks entails a lot more than stacking up layers. Several topics require practical and theoretical knowledge beyond this course:

- Weight initialization
- Regularization (e.g., via dropout)
- Which non-linearities to use
- Which loss functions to use
- How to monitor and adjust the learning process (e.g., optimizers and learning rates) to avoid dying neurons and overfitting



# Notes

# Notes

## Extras

# Full derivation for logistic regression

- First, we need some notable derivatives:

$$\frac{\partial \log(x)}{\partial x} = \frac{1}{x}$$

$$\frac{\partial \sigma(x)}{\partial x} = \sigma(x)(1 - \sigma(x))$$

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f}{\partial g} \cdot \frac{\partial g}{\partial x} \rightarrow \text{chain rule}$$

# Full derivation for logistic regression

- Then:

$$\begin{aligned}\frac{\partial \mathcal{L}_x(\mathbf{w})}{\partial w_j} &= -\partial [y \log \sigma(\mathbf{w}\mathbf{x}) + (1 - y) \log(1 - \sigma(\mathbf{w}\mathbf{x}))] \\ &= -[\partial y \log \sigma(\mathbf{w}\mathbf{x}) + \partial(1 - y) \log(1 - \sigma(\mathbf{w}\mathbf{x}))] \\ &= -\frac{y}{\sigma(\mathbf{w}\mathbf{x})} \partial \sigma(\mathbf{w}\mathbf{x}) - \frac{1 - y}{1 - \sigma(\mathbf{w}\mathbf{x})} \partial(1 - \sigma(\mathbf{w}\mathbf{x})) \rightarrow \text{chain rule} \\ &= -\left[ \frac{y}{\sigma(\mathbf{w}\mathbf{x})} - \frac{1 - y}{1 - \sigma(\mathbf{w}\mathbf{x})} \right] \partial \sigma(\mathbf{w}\mathbf{x}) \rightarrow \text{re-arrange}\end{aligned}$$

- *Exercise: plug-in the derivative of the Sigmoid and re-arrange yourself to reach:*

$$\dots = [\sigma(\mathbf{w}\mathbf{x} - y)] x_j$$

# Full derivation for logistic regression

- In case you were wondering:

$$\begin{aligned}\frac{\partial \sigma(x)}{\partial x} &= \partial \frac{1}{1 + e^{-x}} \\ &= \partial [1 + e^{-x}]^{-1} \\ &= \frac{e^{-x}}{1 + e^{-x}} \frac{1}{1 + e^{-x}} \\ &= \frac{(1 + e^{-x}) - 1}{1 + e^{-x}} \sigma(x) \\ &= \sigma(x)(1 - \sigma(x))\end{aligned}$$

- *Exercise, derive:*

$$\frac{\partial \log \sigma(x)}{\partial x} = \sigma(-x)$$

# Why MSE and cross-entropy?


- It turns out that, given some standard assumptions on our models, using those two losses corresponds to doing Maximum Likelihood Estimation. See <https://www.expunctis.com/2019/01/27/Loss-functions.html>.
- If you are curious about the information theory underpinning cross-entropy, read this: <http://colah.github.io/posts/2015-09-Visual-Information>.

# NN activation functions

Several non-linear activation functions have been proposed. A good default options is ReLU.


logistic ("sigmoid")

$$f(z) = \frac{1}{1 + \exp(-z)}.$$




tanh

$$f(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}},$$




hard tanh

$$\text{HardTanh}(x) = \begin{cases} -1 & \text{if } x < -1 \\ x & \text{if } -1 \leq x \leq 1 \\ 1 & \text{if } x > 1 \end{cases}$$



ReLU (Rectified Linear Unit)

$$\text{rect}(z) = \max(z, 0)$$




*Credit: Stanford CS224N.*




# NN regularization via dropout

Dropout's idea is to mask a random set of neuron connections at training time, in order to compel the network to learn redundant paths and avoid overfitting.



(a) Standard Neural Net



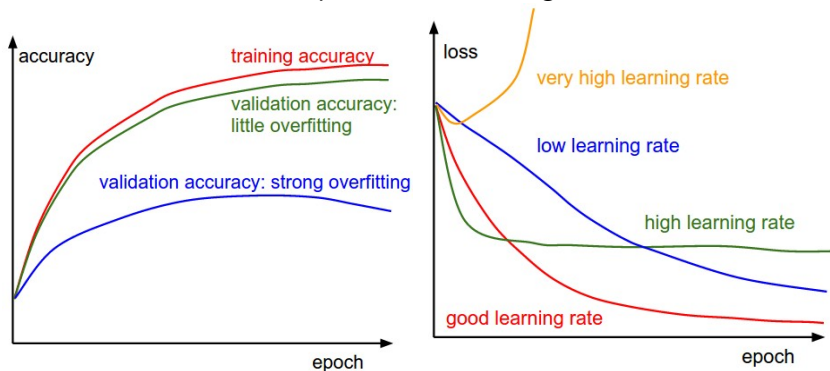
(b) After applying dropout.

*Credit: Srivastava et al.*

*<https://www.cs.toronto.edu/~hinton/absps/JMLRdropout.pdf>.*

# NN under/overfitting and learning rates

Two illustrations on how to spot correct learning behaviour.



*Credit: Andrej Karpathy via Stanford's CS231N.*