

Sentiment Analysis

Jelke Bloem & Giovanni Colavizza

Text Mining
Amsterdam University College

April 25, 2023

Announcements

- Tuesday 02/05: Project update 1
- Tuesday 02/05: Project working session **cancelled**
- Friday 05/05: Liberation day (no lab)
- Monday 15/05: Project update 2
- Monday 15/05: Reading assignment - Bias in word embeddings

BERT reading assignment: Questions

- How well does it work on highly complex human-like tasks, such as common sense of language, slang, and situational jokes?
- If I understood correctly the CLS token represents the whole 'sentence' (text input). However, the text input, the sentence, can be a long text. It doesn't have to be a sentence as in a linguistic sentence. So, can the CLS token represents a whole paper of a few pages?
- Are there any other use cases of BERT that aren't described here? Or does this generally cover everything?
- Does BERT work better on languages that read in a different direction than English?
- Is it possible to not be able to use BERT for certain language tasks?
 - ▶ Sequence length limitation (512 tokens by default)
 - ▶ Generating text (in sequence)
 - ▶ Predicting negative descriptors: A sheepdog is not ... ¹

¹What BERT is not: Lessons from a new suite of psycholinguistic diagnostics for language models <https://arxiv.org/abs/1907.13528v1>

BERT reading assignment: Questions

- Could BERT be made more efficient with more pre-training?
- Could BERT be used without fine-tuning?
- Has GPT-3 changed its fundamental architecture from the GPT discussed in the paper, and if not, how did it manage to outperform BERT as a chatbot (which includes questions and answers, something at which the paper considers it limited)?
- We see that OpenAI's (Chat)GPT model has inherent biases in deployment because of the data it was trained on. Does the same go for BERT? Assuming it does, how would such a model be used in deployment when we can't pinpoint at what stage the biases were infused into the model due to its black box structure?
- Knowing that machine learning models tend to exhibit bias, how does BERT compare to other language models in this regard?
- Models can be highly performant and give good fake results. How can we create more truthful models or even check for it?

BERT reading assignment: Questions

- The authors did not really discuss the prospects of future research using BERT, so what could be potential implications of such embedded architectures?
- How possible is it to install/download these kinds of models on your own computer and run them yourself, since the code and pre-trained models are public? Could you show us how to go about this?
 - ▶ Training only feasible for small datasets or by renting powerful AWS instances. Pre-training BERT would cost about \$4000
 - ▶ Fine-tuning BERT-base: GPU with 12GB of RAM
 - ▶ Inference possible on any system, just slower
 - ▶ Some accessible code in my colleague Jelle Zuidema's course:
https://clclab.github.io/ANCM/05_extra_lab_ICCN-EyeTracking.html

Sentiment Analysis

★★★★★ 6 maanden geleden

This is really a top notch sea and honestly my favorite sea out of all of them. I learned that the water gets saltier the further east you go in the sea which makes people more buoyant when swimming. If I could give it 10 stars, I totally would!

Sentiment Analysis

This review is from: Black Dragon T-Shirt 100% Cotton Short Sleeve Shirt Pre-Shrunk (Apparel)

This is, without a doubt, the best black shirt with an angry monochrome dragon perched on two natural pillars on a cliff that I have ever seen. I know that when I get married, this is the undershirt I'll wear. The amount of awesome displayed on your chest canvas while wearing this shirt, obviously a shirt given to man by Zeus himself, is currently impossible to calculate using our current mathematical constructs. We actually need to devise a new form of mathematics which we should call Wurm Theory in order to parse the data.

Demo

<https://text2data.com/Demo>

Definition

- Sentiment Analysis: The study of **opinions, attitudes and emotions** towards an entity expressed in a text
 - ▶ Persons, objects, events, topics
- Text classification of subjective information
- A one-dimensional scale of opinions is often assumed
 - ▶ like/dislike
 - ▶ for/against
 - ▶ positive/negative
 - ▶ five stars - one star
- Negative and positive polarity
 - ▶ Negative polarity item

Tasks

- Sentiment Classification
 - ▶ Classify the opinions expressed in a text or document into categories or on a scale
- Subjectivity Classification
 - ▶ Detect to what extent a text contains opinions, evaluations, emotions, beliefs etc.
- Opinion Summarization
 - ▶ Summarize the opinions generated by describing a specific entity
- Opinion Retrieval
 - ▶ Retrieve documents that express an opinion about an entity of interest
- Emotion Detection
- Sarcasm/Irony Detection
- Opinion Spam Detection
 - ▶ Messages/reviews that aim to distort the public perception of an entity

Opinions on entities?

- ① I bought an iPhone a few days ago.
 - ▶ a fact, not an opinion
- ② It was such a nice phone.
 - ▶ positive opinion
 - ▶ on the phone as a whole
 - ▶ from the author of the review
- ③ The touch screen was really cool.
 - ▶ positive opinion
 - ▶ on a component of the phone
 - ▶ from the author of the review
- ④ The voice quality was clear too.
 - ▶ positive opinion
 - ▶ on a characteristic of the phone
 - ▶ from the author of the review

Opinions on entities?²

- ⑤ Although the battery life was not long, that is ok for me.
 - ▶ negative opinion?
 - ▶ on a characteristic of the phone
 - ▶ from the author of the review
- ⑥ However, my mother was mad with me as I did not tell her before I bought it.
 - ▶ negative emotion/opinion
 - ▶ towards the author of the review
 - ▶ from the mother of the author of the review
- ⑦ She also thought the phone was too expensive, and wanted me to return it to the shop.
 - ▶ negative opinion
 - ▶ on a characteristic of the phone
 - ▶ from the mother of the author of the review

²Example from Liu (2010) Sentiment Analysis and Subjectivity. In Indurkha & Damerau (eds.) Handbook of Natural Language Processing, CRC Press: Chapter 26.

Entities

- Entity of interest can be a product, person, event, organization or topic
- Entities can have components (battery, screen)
- Entities can have attributes (size, weight, price)
- Or just 'object features'

Opinions

- A view, attitude, emotion, or appraisal on a feature from an opinion holder
- **polarity**: positive, negative or neutral orientation
- **emotion**: subjective feelings and thoughts. Many categorization schemes exist, e.g. VADC, Valence - Arousal - Dominance - Concreteness
- **opinion holder**: a person or organization. Often the author of a post, or sometimes another party that is mentioned (e.g. in a news article)

Unit of analysis

Sentiment of...

- A community
- A specific other person
- The user/author
- Expressed in a document, e.g. a song
- Clause, sentence or paragraph
- Entity or object feature

Applications

- Movies: is this review positive or negative?
- Product: what do people on Twitter think about the new iThingy?
- Advertising: show an ad of our product to people who expressed positive sentiment on a similar product
- Politics: how much do people like this political party currently?
- Economics: is the Twitter stock going to go up or down?
- Users: what other users share similar feelings towards an entity of interest?
- Public sentiment: how is consumer confidence?
- Customer support: is this customer happy or unhappy?
- ...

Characteristics

Sentiment analysis is a type of text classification, but...

- Relatively few categories
- Categories not independent, but e.g. ordinal
- Often not connected to particular domains, topics or users
- Opinions are expressed in complex ways
 - ▶ Lexical content may not be enough
- Ordering effects
- Use of rhetorical devices such as sarcasm and irony
- Information often left implicit
- Different sentiments expressed within a text or even sentence
- Relevance of smileys/emoji

Characteristics

- Negation is complicated
- What is the scope?

*I do **not** call this film a comedy movie and it's bad*

- Scope ambiguity

I wasn't watching all the time

= I was watching none of the time, or I was watching some of the time?

- Long distance negation

*I really like horror movies, but I do **not** feel the same about this one*

Methods

- Supervised classification
 - ▶ e.g. (binarized) Multinomial Naive Bayes
 - ▶ Many datasets available, e.g. movie review polarity datasets, Amazon reviews
 - ▶ Plenty of self-annotated data
- Use of lexical resources
 - ▶ SentiWordNet
 - ▶ WordNet synsets with sentiment scores: positivity, negativity and objectivity
 - ▶ Whissell's Dictionary of Affective Language
 - ▶ About 9000 words rated in terms of their Pleasantness, Activation, and Imagery (concreteness)

Methods

- Use of lexical resources
 - ▶ Regular WordNet (Hu & Liu, 2004)
 - ▶ Begin with known adjectives: “great”, “excellent”, “awful” etc
 - ▶ For unknown adjectives, measure distance to seed adjectives
 - ▶ Use k-nearest neighbour style classification to decide
- Hu & Liu, 2004 developed their own polarity lexicon using this method (2006 positive and 4783 negative words)

Methods

- Supervised classification

Useful features:

- Terms and their (weighted) frequency
- Part-of-speech tags
- Opinion words and phrases (“what a trainwreck!”)
- Negation
- Emoticons/emoji
- Syntactic dependencies

Evaluation issues

- Low inter-rater reliability, hard to get a ground truth
- STS-Gold dataset: Full agreement only on 2/3rd of tweets (for negative/positive/neutral)
- Krippendorff's alpha of 0.765 for tweet-level annotation
- 0.416 for entity-level annotation per tweet

Sentiment is subjective and can be reader-dependent